

Chapter 6 – Procedures

The 80x86 architecture enables implementation of procedures that are similar to those in a high-level language. Procedures use the hardware stack for several purposes. This chapter begins with a discussion of the 80x86 stack and then turns to important procedure concepts-how to call a procedure and return from one, parameter passing, local data, and recursion. The concluding section describes how procedures are implemented in one architecture that does not have a hardware stack.

Chapter 6 – Procedures

6.1 The 80x86 Stack

Programs in this book have allocated stacks with the code

```
.STACK 4096
```

This **.STACK** directive tells the assembler to reserve 4096 bytes of uninitialized storage. The operating system initializes ESP to the address of the first byte above the 4096 bytes in the stack. A larger or smaller stack could be allocated, depending on the anticipated usage in the program.

Chapter 6 – Procedures

6.1 The 80x86 Stack

The stack is most often used by **pushing** words or doublewords on it or by **popping** them off it. This pushing or popping is done automatically as part of the execution of call and return instructions ([Section 6.2](#)). It is also done manually with push and pop instructions.

Source code for a **push instruction** has the syntax

push *source*

The *source* operand can be a register 16, a register 32, a segment register, a word in memory, a doubleword in memory, an immediate byte, an immediate word, or an immediate doubleword.

The only byte-size operand is immediate, and as you will see, multiple bytes are pushed on the stack for an immediate byte operand.

Chapter 6 – Procedures

6.1 The 80x86 Stack

[Figure 6.1](#) lists the allowable operand types. The usual mnemonic for a push instruction is just push. However, if there is ambiguity about the size of the operand (as would be with a small immediate value), then you can use pushw or pushd mnemonics to specify word or doublewordsize operands, respectively.

Operand	Clock Cycles			Number of Bytes	opcode
	386	486	Pentium		
Register	2	1	1	1	
EAX or AX					50
ECX or CX					51
EDX or DX					52
EBX or BX					53
ESP or SP					54
EBP or BP					55
ESI or SI					56
EDI or DI					57

Chapter 6 – Procesures

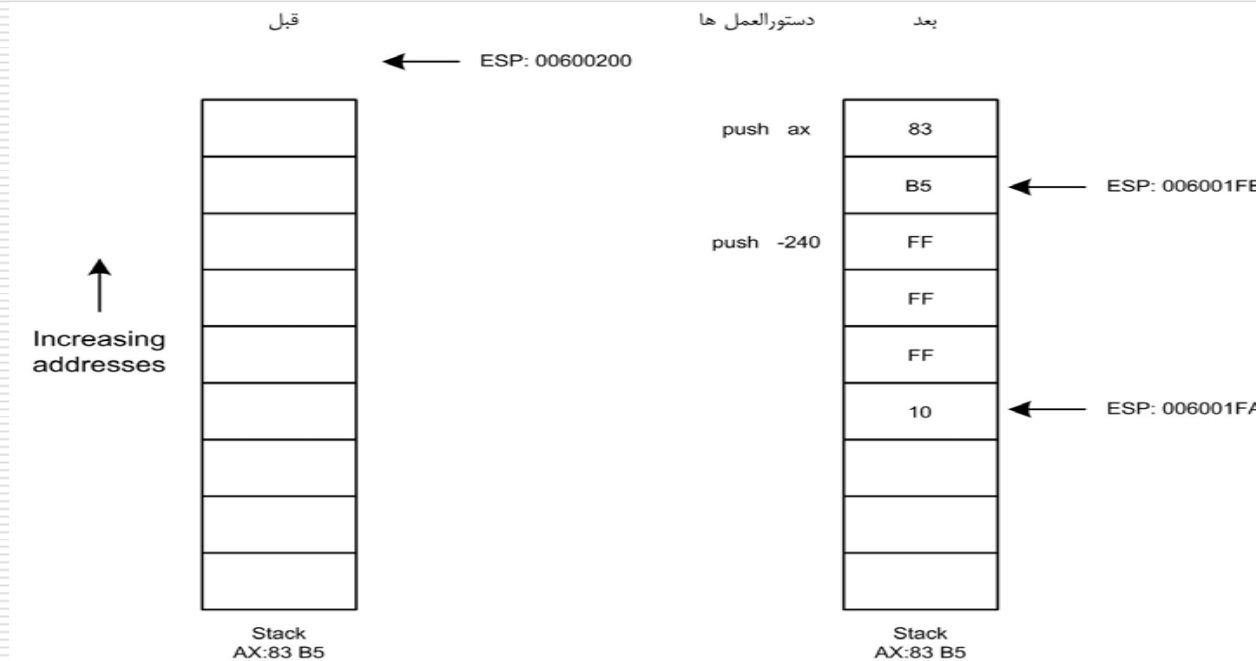
6.1 The 80x86 Stack

Operand	Clock Cycles			Number of Bytes	opcode
	386	486	Pentium		
segment register	2	3	1		
CS				1	0E
DS				1	1E
ES				1	06
SS				1	16
FS				2	0F A0
GS				2	0F A8
memory word	5	4	2	2+	FF
memory doubleword	5	4	2	2+	FF
Immediate byte	2	1	1	2	6A
Immediate word	2	1	1	32+	68
Immediate doubleword	2	1	1	5	68

Chapter 6 – Procedures

6.1 The 80x86 Stack

Example: We now show an example of execution of two push instructions. It assumes that ESP initially contains 00600200. The first push instruction decrements ESP to 006001FE and then stores the contents of AX at that address. Notice that the low and high-order byte are reversed in memory. The second push decrements ESP to 006001FA and stores FFFFFFF10 (-240_{10}) at that address.



Chapter 6 – Procedures

6.1 The 80x86 Stack

As additional operands are pushed onto the stack, ESP is decremented further and the new values are stored. [No push instruction affects any flag bit.](#)

Notice that a stack "grows downward," contrary to the image that you may have of a typical software stack. Also notice that the only value on the stack that is readily available is the last one pushed; it is at the address in ESP. Furthermore, ESP changes frequently as you push values and as procedure calls are made. In [Section 6.3](#) you will learn a way to establish a fixed reference point in the middle of the stack using the EBP register, so that values near that point can be accessed without having to pop off all the intermediate values.

Chapter 6 – Procedures

6.1 The 80x86 Stack

Pop instructions do the opposite job of push instructions.

Each pop instruction has the format

pop *destination*

where *destination* can reference a word or doubleword in memory, any register 16, any register 32, or any segment register except CS. (The push instruction does not exclude CS.) The pop instruction gets a word-size value from the stack by copying the word at the address in ESP to the destination, then incrementing ESP by 2. Operation for a doubleword value is similar, except that ESP is incremented by 4.

Chapter 6 – Procedures

6.1 The 80x86 Stack

[Figure 6.2](#) gives information about pop instructions for different destination operands.

Operand	Clock Cycles			Number of Bytes	opcode
	386	486	Pentium		
Register	4	1	1	1	
EAX or AX					58
ECX or CX					59
EDX or DX					5A
EBX or BX					5B
ESP or SP					5C
EBP or BP					5D
ESI or SI					5E
EDI or DI					5F

Chapter 6 – Procedures

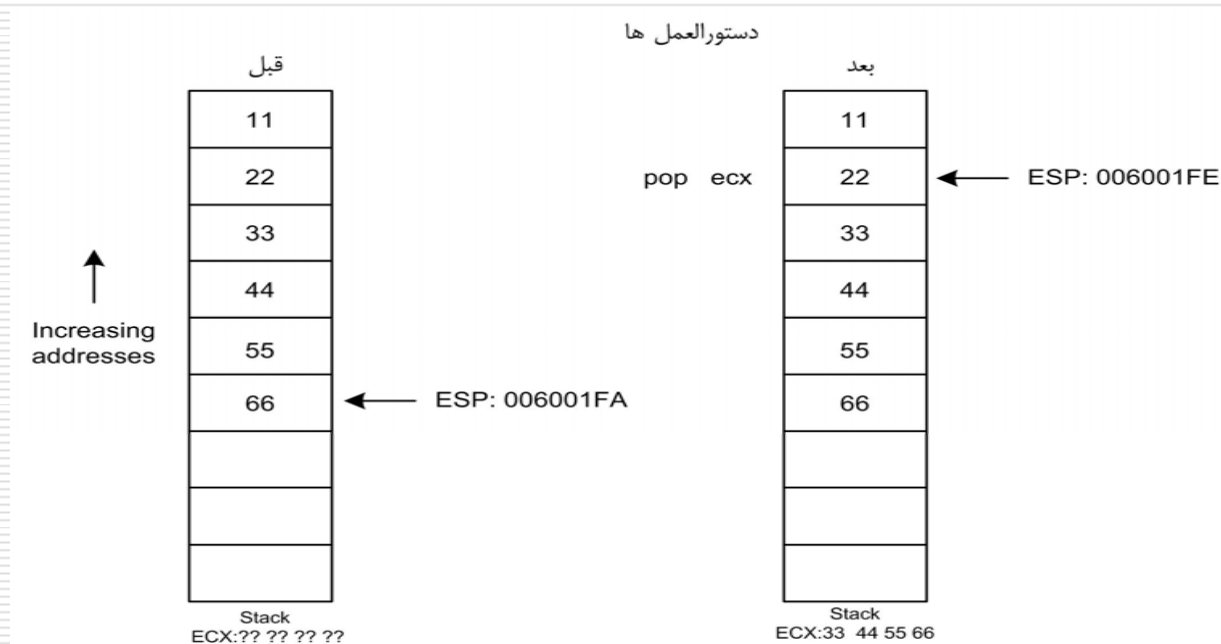
6.1 The 80x86 Stack

Operand	Clock Cycles			Number of Bytes	opcode
	386	486	Pentium		
segment register	7	3	3		
DS				1	1F
ES				1	07
SS				1	17
FS				1	0F A1
GS				2	0F A9
memory word	5	6	3	2+	8F
memory doubleword	5	6	3	2+	8F

Chapter 6 – Procedures

6.1 The 80x86 Stack

This example shows how pop instructions work. The doubleword at the address in ESP is copied to ECX before ESP is incremented by 4. The values popped from the stack are physically still there even though they logically have been removed. Note again that the bytes of a doubleword are stored backward in memory in the 80x86 architecture, but forward in the ECX register.



Chapter 6 – Procedures

6.1 The 80x86 Stack

One use of push and pop instructions is to save the contents of a register temporarily on the stack. We have noted previously that registers are a scarce resource when programming. Suppose, for example, that you are using EDX to store some program variable but need to do a division that requires you to extend a dividend into EDX:EAX prior to the operation. One way to avoid losing your value in EDX is to push it on the stack.

push	edx	;save variable
cdq		;extend dividend to quadword
idiv	Divisor	;divide
pop	edx	;restore variable

This example assumes that you don't need the remainder the division operation puts in EDX. If you do need the remainder, it could be copied somewhere else before popping the value stored on the stack back to EDX.

As the above example shows, push and pop instructions are often used in pairs.

Chapter 6 – Procedures

6.1 The 80x86 Stack

In addition to the ordinary push and pop instructions, special mnemonics push and pop flag registers. These mnemonics are pushf (pushfd for the extended flag register) and popf (popfd for the extended flag register). These are summarized in [Fig. 6.3](#). They are often used in procedure code. Obviously popf and popfd instructions change flag values; these are the only push or pop instructions that change flags.

Operand	Clock Cycles			Number of Bytes	opcode
	386	486	Pentium		
pushf Pushfd	4	4	3	1	9C
popf popfd	5	9	4	1	9D

Chapter 6 – Procedures

6.1 The 80x86 Stack

The 80x86 architecture has instructions that push or pop all general purpose registers with a single instruction. The pushad instruction pushes EAX, ECX, EDX, EBX, ESP, EBP, ESI and EDI, in this order. The value pushed for ESP is the address it contains before any of the registers are pushed. The popad instruction pops the same registers in the opposite order, except that the value for ESP is discarded. Popping the registers in the reverse order ensures that if these instructions are used in a pushad-popad pair, each register (except ESP) will get back its original value. [Figure 6.4](#) shows the push all and pop all instructions, including the pusha and popa instructions that push and pop the 16-bit registers.

Operand	Clock Cycles			Number of Bytes	opcode
	386	486	Pentium		
pusha pushad	18	11	5	1	60
popa popad	24	9	5	1	61

Chapter 6 – Procedures

6.1 The 80x86 Stack

Finally, a note of caution. Although the Intel architecture allows 16-bit or 32-bit quantities to be pushed on the stack, some operating systems (including Microsoft Windows NT) require parameters used in system calls to be on doubleword boundaries, that is, a parameter's address must be a multiple of 4. The stack starts on a doubleword boundary, but to maintain this alignment, only doublewords should be pushed on the stack prior to a system call.

Chapter 6 – Procedures

6.2 Procedure Body, Call and Return

The word ***procedure*** is used in high-level languages to describe a subprogram that is almost a self-contained unit. The main program or another subprogram can call a procedure by including a statement that consists of the procedure name followed by a parenthesized list of arguments to be associated with the procedure's formal parameters.

In assembly language and in some high-level languages the term *procedure* is used to describe both types of subprograms, those that return values and those that do not.

Procedures are valuable in assembly language for the same reasons as in high-level languages. They help divide programs into manageable tasks and they isolate code that can be used multiple times within a single program or that can be saved and reused in several programs.

Chapter 6 – Procedures

6.2 Procedure Body, Call and Return

This section describes how to write 80x86 procedures, as well as how to assemble and link them using Microsoft software. Information is included on [how to define a procedure](#), and [how to transfer execution control to a procedure and back to the calling program](#). We show how the stack is used to save register contents, so that a procedure returns to the caller with almost all registers unchanged. Other important concepts to be considered with procedures are [how to pass arguments to a procedure](#) and [how to implement local variables in a procedure body](#).

The code for a procedure always follows a `.CODE` directive. The body of each procedure is bracketed by two directives, `PROC` and `ENDP`. Each of these directives has a label that gives the name of the procedure. With the Microsoft Macro Assembler, the `PROC` directive allows several attributes to be specified; we are only going to use one, `NEAR32`. This attribute says that the procedure will be located in the same code segment as the calling code and that 32-bit addresses are being used. These choices are normal for flat 32-bit memory model programming.

Chapter 6 – Procedures

6.2 Procedure Body, Call and Return

```
; procedure structure example  
; Author: R. Detmer  
; Date: revised 10/97
```

```
.386
```

```
.MODEL FLAT
```

```
ExitProcess PROTO NEAR32 stdcall, dwExitCode:DWORD
```

```
.STACK 4096 ; reserve 4096-byte stack
```

```
.DATA ; reserve storage for data
```

```
Count1 DWORD 11111111h
```

```
Count2 DWORD 22222222h
```

```
Total1 DWORD 33333333h
```

```
Total2 DWORD 44444444h
```

```
; other data here
```

Chapter 6 – Procedures

6.2 Procedure Body, Call and Return

```
.CODE                                ; program code

Initialize      PROC      NEAR32
                mov       Count1,0      ;zero first count
                mov       Count2,0      ;zero second count
                mov       Total1,0      ;zero first total
                mov       Total2,0      ;zero second total
                mov       ebx,0         ;zero balance
                ret         ;return
Initialize      ENDP

_start:                                ;program entry point
                call      Initialize    ;initialize variables
; -- other program tasks here

                call      Initialize    ;reinitialize variables
; -- more program tasks here
                INVOKE     ExitProcess,0 ;exit with return code 0
PUBLIC _start    ;make entry point public
END              ;end of source code
```

Chapter 6 – Procedures

6.2 Procedure Body, Call and Return

In [Fig. 6.5](#) the procedure *Initialize* is bracketed by PROC and ENDP. The distance attribute NEAR32 declares this to be a near procedure. Although this example shows the procedure body prior to the main code, it could also have been placed afterwards. Recall that execution of a program does not necessarily begin at the first statement of the code segment; the statement identified by the label `_start` marks the first instruction to be executed.

Most of the statements of procedure *Initialize* are ordinary mov instructions. These could have been used in the main program at the two places that the call statements are coded; however, using the procedure makes the main code both shorter and clearer. The procedure affects doublewords defined in the program's data segment and the EBX register; it has no local variables.

Chapter 6 – Procedures

6.2 Procedure Body, Call and Return

When the main program executes, the instruction

call Initialize

transfers control from the main code to the procedure. The main program calls the procedure twice; in general, a procedure may be called any number of times. The return instruction

ret

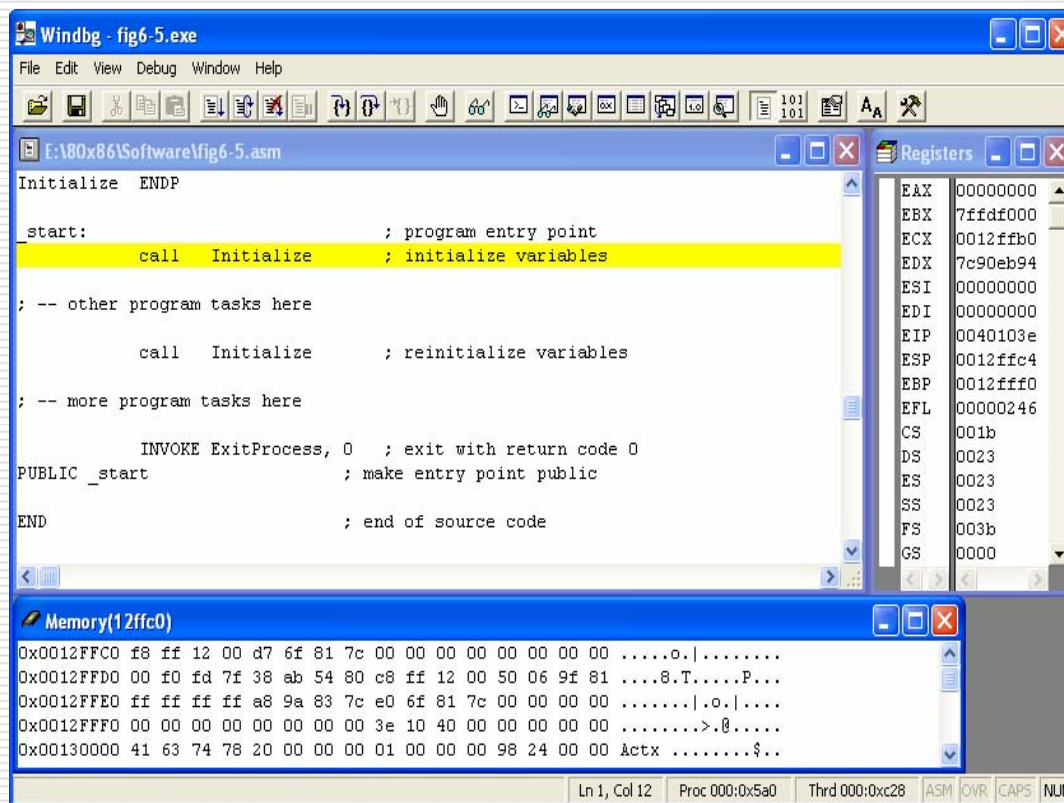
transfers control from the procedure back to the caller; **there is almost always at least one ret instruction in a procedure and there can be more than one.** If there is only one ret, it is ordinarily the last instruction in the procedure since subsequent instructions would be unreachable without "spaghetti code." Although a call instruction must identify its destination, the ret does not. control will transfer to the instruction following the most recent call. The 80x86 uses the stack to store the return address.

Chapter 6 – Procedures

6.2 Procedure Body, Call and Return

When the example program in [Fig. 6.5](#) is assembled, linked, and executed, there is no visible output. However, it is informative to trace execution with a tool like WinDbg.

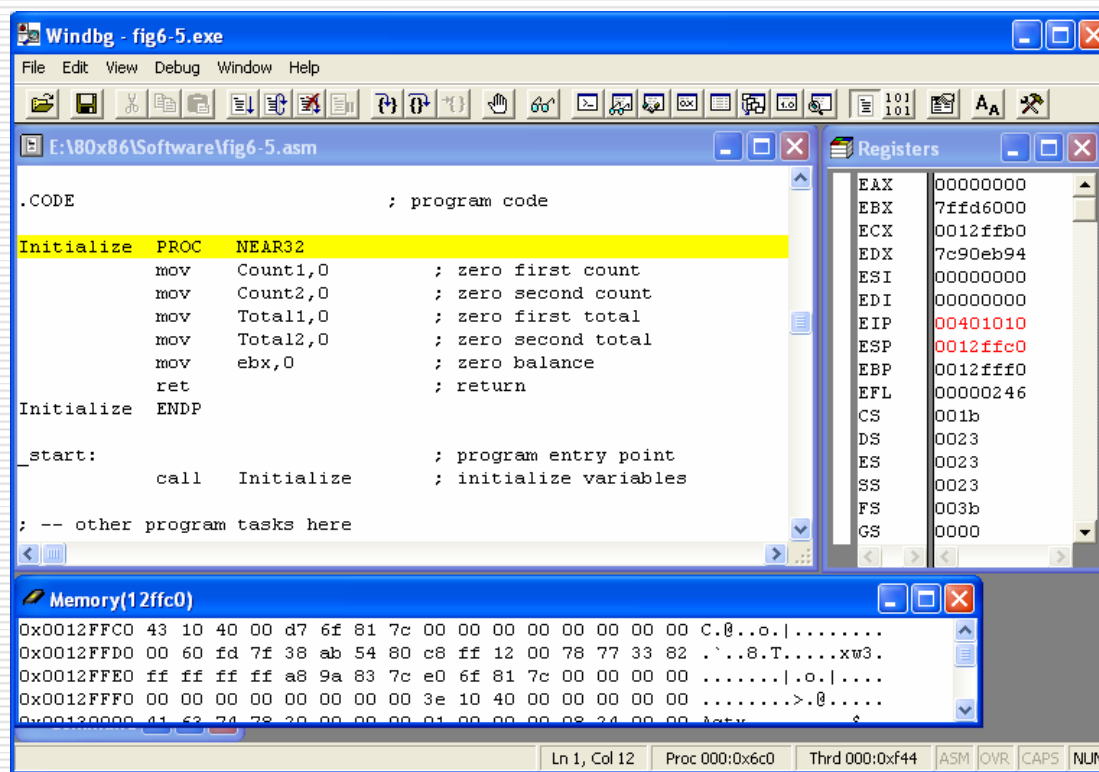
[Figure 6.6](#) show WinDbg's initial display. Note that ESP contains 0012FFC4. The memory window has been opened to start at address 0012FFC0, 12 bytes down into the stack. The EIP register contains 0040103E, the address of the first instruction to be executed (the first call).



Chapter 6 – Procedures

6.2 Procedure Body, Call and Return

Figure 6.7 shows the new state after this statement is executed. The EIP register now contains 00401010, the address of the first statement in procedure *Initialize*. The ESP register contains 0012FFE0, so four bytes have been pushed onto the stack. Looking in memory at this address, you see 43 10 40 00—that is, 00401043, an address five bytes larger than the address of the first call. If you examine the listing file for the program, you see that the each call instruction takes five bytes of object code, so that 00401043 is the address of the instruction following the first call.



Chapter 6 – Procedures

6.2 Procedure Body, Call and Return

In general, a call instruction pushes the address of the next instruction (the one immediately following the call) onto the stack and then transfers control to the procedure code. A near call instruction works by pushing the EIP to the stack and then changing EIP to contain the address of the first instruction of the procedure.

Return from a procedure is accomplished by reversing the above steps. A ret instruction pops EIP, so that the next instruction to be executed is the one at the address that was pushed on the stack.

Recall that 80x86 programming can be done using either a flat memory model or a segmented memory model. With a segmented memory model a procedure may be in a different segment from the calling code. In fact, with 16-bit segmented programming, segments were limited to 65,536 bytes, so procedures were often in separate segments. The 80x86 architecture uses a far call to transfer control to a procedure in a different memory segment: A far call pushes both EIP and CS onto the stack. A far return pops both off the stack. With 32-bit flat memory model programming, there is no need for anything other than near calls.

Chapter 6 – Procedures

6.2 Procedure Body, Call and Return

The syntax of the 80x86 call statement is

call destination

[Figure 6.8](#) lists some of the available 80x86 call instructions, omitting 16-bit forms and forms used primarily for systems programming. No call instruction modifies any flag.

Operand	Clock Cycles			Number of Bytes	opcode
	386	486	Pentium		
Near relative	7+	3	1	5	E8
Near indirect					FF
using register	7+	5	2	2	
using memory	10+	5	2	2+	
far direct	17+	18	4	7	9A
far indirect	22+	17	5	6	FF

Chapter 6 – Procedures

6.2 Procedure Body, Call and Return

For a near relative procedure, the assembler calculates a 32-bit displacement to the destination, and the E8 opcode plus this displacement comprise the five bytes of the instruction. The transfer of control when a procedure is called is similar to the transfer of a relative jump, except that the old contents of EIP are pushed, of course.

Near indirect calls encode a register 32 or a reference to a doubleword in memory. When the call is executed, the contents of that register or doubleword are used as the address of the procedure.

All far calls must provide both new CS contents and new EIP contents. With far direct calls, both of these are coded in the instruction, adding six bytes to the opcode. With far indirect calls, these are located at a six-byte block in memory, and the address of that block is coded in the instruction.

Chapter 6 – Procedures

6.2 Procedure Body, Call and Return

The **return instruction** `ret` is used to transfer control from a procedure body back to the calling point. Its basic operation is simple; it pops the address previously stored on the stack and loads it into the instruction pointer EIP.

A near return just has to restore EIP. A far return instruction reverses the steps of a far call, restoring both EIP and CS; both of these values are popped from the stack.

There are two formats for the `ret` instruction. The more common form has no operand and is simply coded

`ret`

An alternative version has a single operand and is coded

`ret count`

The operand *count* is added to the contents of ESP after completion of the other steps of the return process (restoring EIP and, for a far procedure, CS).

Chapter 6 – Procedures

6.2 Procedure Body, Call and Return

Figure 6.9 ret instruction

Operand	Operand	Clock Cycles			Number of Bytes	Opcode
		386	486	Pentium		
near	none	10+	5	2	1	C3
near	immediate	10+	5	3	3	C2
far	none	18+	13	4	1	CB
far	immediate	18+	14	4	3	CA

If a procedure's PROC directive has the operand NEAR32, then the assembler generates near calls to the procedure and near returns to exit from it. The Microsoft Macro Assembler also has [retn](#) (return near) and [retf](#) (return far) mnemonics to force near or far returns.

Chapter 6 – Procedures

6.2 Procedure Body, Call and Return

procedures and calling programs in separate files

To construct building blocks for large programs, it is often desirable to assemble a procedure or group of procedures separately from the code that calls them. There are a few additional steps required to do this.

- ✓ First, you must assemble the procedures so that their names are visible outside the file containing them.
- ✓ Second, you must let the calling program know necessary information about the external procedures.
- ✓ Finally, you must link the additional .OBJ files to get an executable program.

The `PUBLIC` directive is used to make procedure names visible outside the file containing them. This is the same directive we have been using to make the symbol `_start` visible. In general, its syntax is

```
PUBLIC      symbol1 [, symbol2]...
```

A file may contain more than one `PUBLIC` directive

Chapter 6 – Procedures

6.2 Procedure Body, Call and Return

The EXTRN directive gives the calling program information about external symbols. It has many options, including

```
EXTRN      symbol1:type [,symbol2:type]
```

A file may contain more than one EXTRN directive.

File containing procedure definitions

```
PUBLIC      Procedure1, Procedure2
```

```
.CODE
```

```
Procedure1  PROC NEAR32
```

```
...
```

```
Procedure1  ENDP
```

```
Procedure2  PROC NEAR32
```

```
...
```

```
Procedure2  ENDP
```

```
END
```

Chapter 6 – Procedures

6.2 Procedure Body, Call and Return

File containing procedure calls

```
EXTRN Procedure1:NEAR32, Procedure2:NEAR32
```

```
...  
.CODE  
...  
call Procedure1  
...  
call Procedure2  
...  
END
```

You assemble each of the above files just as if it were the main program. Each assembly produces a .OBJ file. To link the files, simply list all .OBJ files in the link command-you already have been linking your programs with the separately assembled file IO.OBJ.

Chapter 6 – Procedures

6.2 Procedure Body, Call and Return

We conclude this section with a procedure that will calculate the integer square root of a positive integer *Nbr*; that is, the largest integer *SqRt* such that $SqRt * SqRt \leq Nbr$. The procedure code is in [Fig. 6.11](#).

Procedure Root implements the following design.

```
Sqrt := 0;
while Sqrt * Sqrt ≤ Nbr loop
    add 1 to Sqrt;
end while;
subtract 1 from Sqrt;
```


Chapter 6 – Procedures

6.2 Procedure Body, Call and Return

```
; procedure to compute integer square root of number Nbr
; Nbr is passed to the procedure in EAX
; The square root SqRt is returned in EAX
; Other registers are unchanged.
; author: R. Detmer revised: 10/97
```

```
Root      PROC      NEAR32
           push      ebx                ;save registers
           push      ecx
           mov        ebx,0             ;SqRt := 0
WhileLE:   mov        ecx,ebx           ;copy SqRt
           imul       ecx,ebx           ;SqRt*SqRt
           cmp        ecx,eax           ;SqRt*SqRt <= Nbr ?

           jnle       EndWhileLE        ;exit if not
           inc        ebx               ;add 1 to SqRt
           jmp        WhileLE           ;repeat
EndWhileLE:
           dec        ebx               ;subtract 1 from SqRt
           mov        eax,ebx           ;return SqRt in AX
           pop        ecx               ;restore registers
           pop        ebx
           ret                        ;return Root ENDP
```

Chapter 6 – Procedures

6.3 Parameters and Local Variables

Using a high-level language, a procedure definition often includes parameters or formal parameters that are associated with arguments or actual parameters when the procedure is called. For the procedure's **(pass-by-value)** parameters, values of the arguments, which may be expressions, are copied to the parameters when the procedure is called, and these values are then referenced in the procedure using their local names, which are the identifiers used to define the parameters.

In-out (pass-by-location or variable) parameters associate a parameter identifier with an argument that is a single variable and can be used to pass a value either to the procedure from the caller or from the procedure back to the caller.

A common technique for passing parameters is discussed in this section. This technique can be used to pass word-size or doubleword-size values for **in** parameters, or addresses of data in the calling program for **in-out** parameters.

Although simple procedures can be written using only registers to pass parameters, most procedures use the stack to pass parameters.

Chapter 6 – Procedures

6.3 Parameters and Local Variables

We start with a simple example to show how the stack is used to pass parameters. Suppose that the job of a NEAR32 procedure *Add2* is to add two doubleword-size integers, returning the sum in EAX. If the calling program passes these parameters by pushing them on the stack, then its code might look like

push	Value1	;first argument value
push	ecx	;second argument value
call	Add2	;call procedure to find sum
add	esp,8	;remove parameters from stack

Before we look at how the parameter values are accessed from the stack, notice how they are removed from the stack following the call. There is no need to pop them off the stack to some destination; we simply add eight to the stack pointer to move ESP above the parameters.

An alternative to adding *n* to the stack pointer in the calling program is to use `ret n` in the procedure, the version of the return instruction that adds *n* to ESP after popping the return address.

Chapter 6 – Procedures

6.3 Parameters and Local Variables

[Figure 6.12](#) shows how the procedure *Add2* can retrieve the two parameter values from the stack. The procedure code uses the **based** addressing mode. In this mode, a memory address is calculated as the sum of the contents of a base register and a displacement built into the instruction.

```
Add2          PROC    NEAR32  ;add two words passed on the stack
                                   ;return the sum in the EAX register


               push    ebp          ;save EBP
               mov     ebp,esp      ;establish stack frame
               mov     eax,[ebp+8]   ;copy second parameter value
               add     eax,[ebp+12]  ;add first parameter value
               pop     ebp          ;restore EBP
               ret              ;return

Add2          ENDP
```

Chapter 6 – Procedures

6.3 Parameters and Local Variables

This method of passing argument values works as follows. Upon entry to the procedure, the stack looks like the left illustration in [Fig. 6.13](#).



Chapter 6 – Procedures

6.3 Parameters and Local Variables

After the procedure's instructions

```
push    ebp                ;save EBP
mov     ebp,esp            ;establish stack frame
```

are executed, the stack looks like the right illustration in [Fig. 6.13](#).

Eight bytes are stored between the address stored in EBP (and also ESP) and the second parameter value. Therefore parameter 2 can be referenced by [bp+8]. The first parameter value is four bytes higher on the stack; its reference is [bp+12]. The code

```
mov     eax,[bp+8]         ;copy second parameter
add     eax,[bp+12]        ;add first parameter
```

uses the values from memory locations in the stack to compute the desired sum.

Chapter 6 – Procedures

6.3 Parameters and Local Variables

You may wonder **why EBP is used at all**. Why not just use ESP as a base register? The principal reason is that ESP is likely to change, but the instruction `mov ebp,esp` loads EBP with a *fixed* reference point in the stack. This fixed reference point will not change as the stack is used for other purposes, for example, to push additional registers or to call other procedures.

Some procedures need to allocate stack space for local variables, and most procedures need to save registers as illustrated in [Fig. 6.11](#). Instructions to accomplish these tasks, along with the instructions

```
push    ebp           ;save EBP
mov     ebp,esp       ;establish stack frame
```

make up the entry code for a procedure. However, [the two instructions here are always the first entry code instructions](#). Because they are, you can count on the last parameter being exactly eight bytes above the reference point stored in EBP. The EBP register itself is always the first pushed and last popped so that upon return to the calling program it has the same value as prior to the call.

Chapter 6 – Procedures

6.3 Parameters and Local Variables

We now show how the stack can provide space for local variables. For this purpose, we revisit the algorithm for computing the greatest common divisor of two integers that appeared in [Programming Exercises 5.3](#).

```
gcd := number1;
remainder := number2;

until (remainder = 0) loop
    dividend := gcd;
    gcd := remainder;
    remainder := dividend mod gcd;
end until;
```

[Figure 6.14](#) shows this design implemented as a NEAR32 procedure that computes the greatest common divisor of two doubleword-size integer values passed to the procedure on the stack, returning the GCD in EAX. [Figure 6.14](#) includes more than the procedure itself. It shows a complete file, ready for separate assembly.

Chapter 6 – Procedures

6.3 Parameters and Local Variables

[Figure 6.14](#) includes more than the procedure itself. It shows a complete file, ready for separate assembly.

PUBLIC GCD

```
GCD      PROC      NEAR32
        push      ebp                ;establish stack frame
        mov       ebp,esp
        sub       esp,4             ;space for one local doubleword
        push      edx                ;save EDX
        pushf     ;save flags
        mov       eax,[ebp+8]        ;get Number1
        mov       [ebp-4],eax        ;GCD := Number1
        mov       edx,[ebp+12]       ;Remainder := Number1
until0:  mov       eax,[ebp-4]        ;Dividend := GCD
        mov       [ebp-4],edx        ;GCD := Remainder
        mov       edx,0              ;extend Dividend to doubleword
        div       DWORD PTR [ebp-4];Remainder in EDX
        cmp       edx, 0             ;remainder = 0?
        jnz       until0            ;repeat if not
```

Chapter 6 – Procedures

6.3 Parameters and Local Variables

```
        mov     eax,[ebp-4]      ;copy GCD to EAX
        popf                    ;restore flags
        pop     edx             ;restore EDX
        mov     esp,ebp         ;restore ESP
        pop     ebp             ;restore EBP
        ret     8               ;return, discarding parameters
GCD     ENDP
        END
```

In this procedure, *gcd* is stored on the stack until it is time to return the value in EAX.
The instruction


```
sub     esp,4      ;space for one local doubleword
```

moves the stack pointer down four bytes, reserving one doubleword of space below where EBP was stored and above where other registers are stored.

Chapter 6 – Procedures

6.3 Parameters and Local Variables

After EDX and the flags register are pushed, the stack has the contents shown in [Fig. 6.15](#). Now the local variable *gcd* can be accessed as `[ebp-4]`, since it is four bytes below the fixed reference point stored in EBP.



The rest of the procedure is a straightforward implementation of the design. In this example, a register could have been used to store *gcd*, but many procedures have too many local variables to store them all in registers. Within reason, you can reserve as many local variables on the stack as you wish, accessing each by `[ebp-offset]`.

Chapter 6 – Procedures

6.3 Parameters and Local Variables

Finally, consider the exit code for the procedure.

```
    popf                ;restore flags
    pop     edx          ;restore EDX
    mov     esp,ebp      ;restore ESP
    pop     ebp          ;restore EBP
    ret     8            ;return, discarding parameters
```

The first two pop instructions simply restore the flag register and EDX; these instructions are popped in the opposite order in which they were pushed. It may seem that the next instruction should be add sp,4 to undo the effects of the corresponding subtraction in the entry code.

However, the instruction mov esp,ebp accomplishes the same objective more efficiently, working no matter how many bytes of local variable space were allocated, and without changing flags like an add instruction.

Finally, we are using the ret instruction with operand 8, so that for this procedure the calling program must *not* remove parameters from the stack; this task is accomplished by the procedure.

Chapter 6 – Procedures

6.3 Parameters and Local Variables

[Figure 6.16](#) summarizes typical entry code and exit code for a procedure. High-level language compilers generate similar code for subprograms.

Entry code:

```
push    ebp                ;establish stack frame
mov     ebp,esp
sub     esp,n              ; n bytes of local variables space
push    ...
...
push    ...
pushf                   ;save flags
```

Exit code

```
popf                   ;restore flags
pop     ...
...
pop     ...
mov     esp,ebp          ;restore ESP if local variables used
pop     ebp              ;restore EBP
ret                    ;return
```

Chapter 6 – Procedures

6.3 Parameters and Local Variables

How can a high-level language implement variable parameters? How can large parameters such as an array, a character string, or a record be efficiently passed to a procedure? Either of these can be implemented by passing the address of an argument rather than the value of the argument to the procedure. [Figure 6.17](#) shows a procedure that might implement the Pascal procedure with header

```
PROCEDURE Minimum(A : IntegerArray;  
                  Count : INTEGER;  
                  VAR Min : INTEGER);  
(* Set Min to smallest value in A[1], A[2], ..., A[Count] *)
```

In this implementation the addresses of arguments corresponding to *A* and *Min* are passed to procedure *Minimum*. The procedure uses register indirect addressing, first to examine each array element, and at the end to store the smallest value.

Chapter 6 – Procedures

6.3 Parameters and Local Variables

Note that since the Count parameter is word-size, the address of the first parameter is 14 bytes above the fixed base point, four bytes for EBP, four bytes for the return address, four bytes for the address of Min, and two bytes for the value of Count.

Calling code for procedure *Minimum* could look like the following.

```
lea    eax,Array        ;Parameter 1: address of Array
push   eax
push   Count            ;Parameter 2: value of Count
lea    eax,Min          ;Parameter 3: address of Min
push   eax
call   Minimum          ;call procedure
add    esp,10           ;discard parameters
```

After this code is executed, the smallest value from Array will be in the word referenced by Min.

Chapter 6 – Procedures

6.3 Parameters and Local Variables

storing local variables in the data segment

It is perfectly legal for a procedure to store local variables in the data segment. The `.DATA` directive can be included in a file used for separate assembly of procedures. In fact, a program may have multiple `.DATA` directives, although this is generally not necessary. You should normally keep variables as local as possible.

Because 80x86 instructions are often the output of a compiler, the 80x86 architecture includes additional instructions to facilitate implementation of procedures. The `enter` instruction has syntax

```
enter      localBytes, nestingLevel
```

When *nestingLevel* is zero, this does precisely the job of the following familiar instructions:

```
push      ebp  
mov       ebp, esp  
sub       esp, localBytes
```


Chapter 6 – Procedures

6.3 Parameters and Local Variables

If *nestingLevel* > 0, the enter instruction also pushes the stack frame pointers from *nestingLevel-1* levels back onto the stack above the new frame pointer. This gives this procedure easy access to the variables of procedures in which it is nested. If used, an enter instruction would normally be the first instruction in a procedure.

The leave instruction reverses the actions of the enter instruction. Specifically, it does the same thing as the instruction pair

mov	esp,ebp	;restore ESP
pop	ebp	;restore EBP

and normally would be used immediately before a return instruction.

Chapter 6 – Procedures

6.3 Parameters and Local Variables

You have observed that each program we write exits with the statement

```
INVOKE      ExitProcess,0           ;exit with return code 0
```

INVOKE is not an instruction—MASM references call it a directive. However, it acts more like a macro. In fact, if the directive `.LISTALL` precedes the above line of code, you see the expansion

```
push        +00000000h  
call        ExitProcess
```

This is clearly a call to procedure *ExitProcess* with a single doubleword parameter with value 0.

Chapter 6 – Procedures

6.4 Recursion

6.4 Recursion


A recursive procedure or function is one that calls itself, either directly or indirectly. The best algorithms for manipulating many data structures are recursive.

It is almost as easy to code a recursive procedure in 80x86 assembly language as it is to code a nonrecursive procedure. If parameters are passed on the stack and local variables are stored on the stack, then each call of the procedure gets new storage allocated for its parameters and local variables.

Chapter 6 – Procedures

6.4 Recursion

This section gives one example of a recursive procedure in 80x86 assembly language. It solves the Towers of Hanoi puzzle, pictured in [Fig. 6.18](#) with four disks. **The object of the puzzle is to move all disks from source spindle A to destination spindle B, one at a time, never placing a larger disk on top of a smaller disk.** Disks can be moved to spindle C, a spare spindle. For instance, if there are only two disks, the small disk can be moved from spindle A to C, the large one can be moved from A to B, and finally the small one can be moved from C to B.



Chapter 6 – Procedures

6.4 Recursion

In general, the Towers of Hanoi puzzle is solved by looking at two cases. If there is only one disk, then the single disk is simply moved from the source spindle to the destination. If the number of disks *NbrDisks* is greater than one, then the top (*NbrDisks*-1) disks are moved to the spare spindle, the largest one is moved to the destination, and finally the (*NbrDisks*-1) smaller disks are moved from the spare spindle to the destination. Each time (*NbrDisks*-1) disks are moved, exactly the same procedure is followed, except that different spindles have the roles of source, destination, and spare. [Figure 6.19](#) expresses the algorithm in pseudocode.

Chapter 6 – Procedures

6.4 Recursion

```
procedure Move(NbrDisks, Source, Destination, Spare);
begin
    if NbrDisks = 1
    then
        display "Move disk from ", Source, " to ", Destination
    else
        Move(NbrDisks - 1, Source, Spare, Destination);
        Move(1, Source, Destination, Spare);
        Move(NbrDisks - 1, Spare, Destination, Source);
    end if;
end procedure Move;
begin {main program}
    prompt for and input Number;
    Move(Number, 'A', 'B', 'C');
end;
```

Chapter 6 – Procedures

6.4 Recursion

[Figure 6.20](#) shows 80x86 code that implements the design. The stack is used to pass parameters to procedure *Move*, which is a NEAR32 procedure referencing the data segment for output only. A standard stack frame is established, and registers used by the procedure are saved and restored. The code is a fairly straightforward translation of the pseudocode design. The operator DWORD PTR is required in the statement

```
cmp      DWORD PTR [bp+14],1
```

so that the assembler knows whether to compare words or byte size operands. Similarly, the pushw mnemonic is used several places so that the assembler knows to push wordsize parameters. Notice that the recursive calls are implemented exactly the same way as the main program call, by pushing four parameters on the stack, calling procedure *Move*, then removing the parameters from the stack. However, in the main program the spindle parameters are constants, stored as the low order part of a word since single bytes cannot be pushed on the 80x86 stack.

Chapter 6 – Procedures

6.5 Other Architectures: Procedures Without a Stack

6.5 Other Architectures: Procedures Without a Stack

Not all computer architectures provide a hardware stack. One can always implement a software stack by setting aside a block of memory, thinking of it as a stack, maintaining the stack top in a variable, and pushing or popping data by copying to or from the stack.

Obviously the stack plays a large role in 80x86 procedure implementation. How can you reasonably implement procedures in an architecture that has no stack?

The S/360 architecture includes sixteen 32-bit *general purpose registers* (GPRs), numbered 0 to 15. Addresses are 24 bits long and an address can be stored in any register. The architecture includes addressing modes comparable to direct, register indirect, and indexed.

A procedure is usually called by loading its address into GPR 15 and then executing a *branch and link* instruction that jumps to the procedure code after copying the address of the next instruction into GPR 14. This makes return easy; simply jump to the address in GPR 14.

Chapter 6 – Procedures

6.5 Other Architectures: Procedures Without a Stack

Parameter passing is more challenging. Normally GPR 1 is used to pass the address of a *parameter address list*. This is a list of 32-bit storage locations (32 bits is called a *word* in the S/360 architecture), the first word containing the address of the first parameter, the second word containing the address of the second parameter, etc. To retrieve a word-size parameter, one must first get its address from the parameter address list, then copy the word at that address.

Since the same general purpose registers are normally used for the same tasks each time a procedure is called, problems may occur if one procedure calls another. For instance, a second procedure call would put the second return address into GPR 14, wiping out the first return address. To avoid this problem, the main program and each procedure allocates a block of storage for a *register save area* and puts its address in GPR 13 prior to a procedure call. The procedure then saves general purpose registers 0-12, 14, and 15 in the register save area of the calling program and GPR 13 in its own register save area. This system is relatively complicated compared to using a stack, but it works well except for recursive procedure calls. Since there is only one register save area per procedure, recursive procedure calls are impossible without modifying the scheme.

Chapter 6 – Procedures

Chapter Summary

This chapter has discussed techniques for implementing procedures in the 80x86 architecture. The stack serves several important purposes in procedure implementation. When a procedure is called, the address of the next instruction is stored on the stack before control transfers to the first instruction of the procedure. A return instruction retrieves this address from the stack in order to transfer control back to the correct point in the calling program. Argument values or their addresses can be pushed onto the stack to pass them to a procedure; when this is done, the base pointer EBP and based addressing provide a convenient mechanism for accessing the values in the procedure. The stack can be used to provide space for a procedure's local variables. The stack is also used to "preserve the environment"-for example, register contents can be pushed onto the stack when a procedure begins and popped off before returning to the calling program so that the calling program does not need to worry about what registers might be altered by the procedure.

Chapter 6 – Procedures

Chapter Summary

Recursive algorithms arise naturally in many computing applications. Recursive procedures are no more difficult than nonrecursive procedures to implement in the 80x86 architecture.

Some computer architectures do not have a hardware stack. Nonrecursive procedures can be implemented using registers to store addresses, and memory to save registers when one procedure calls another.